

## Microservices Interview Questions

This will include:

-  Core concepts
  -  Design patterns
  -  Communication
  -  Advanced scenarios
  -  Real-world architecture questions
  -  Behavioral & experience-based questions
- 

### 1 Microservices Basics

#### Q1. What are microservices?

- Independent, small, loosely coupled services.
- Each service runs in its own process, handles a single business capability, and communicates over lightweight protocols (HTTP/REST, gRPC, Kafka).

#### Q2. Monolith vs Microservices

Feature	Monolith	Microservices
Deployment	Single unit	Independent services
Scaling	Whole app	Per service
Tech Stack	Usually one	Can be mixed
Failure	Whole system down	Isolated failures

#### Q3. Why use microservices?

- Scalability, flexibility, faster deployments, fault isolation, tech diversity.

#### Q4. What are the main challenges?

- Distributed transactions
  - Data consistency
  - Communication overhead
  - Service discovery and tracing
  - DevOps complexity
- 

### 2 Architecture & Design

#### Q5. What are the core components of a microservices architecture?

- API Gateway
- Service Registry
- Config Server
- Message Broker
- Centralized Logging & Monitoring
- Circuit Breakers / Resilience Layer
- Security Layer (OAuth2/JWT)

#### Q6. What are design principles of microservices?

- Single Responsibility

- Loose Coupling
- High Cohesion
- Resilience
- Scalability
- Automation (CI/CD)

#### **Q7. Explain domain-driven design (DDD) in microservices**

- Divides system into **bounded contexts**.
- Each microservice represents one bounded context and owns its data and logic.

#### **Q8. What are common design patterns in microservices?**

<b>Pattern</b>	<b>Purpose</b>
API Gateway	Entry point for all requests
Circuit Breaker	Prevents cascading failures
Service Discovery	Auto-registration & lookup
Saga	Distributed transaction handling
CQRS	Separate read and write models
Strangler Fig	Gradual migration from monolith
Bulkhead	Resource isolation

### **3 Communication Between Services**

#### **Q9. How do microservices communicate?**

- **Synchronous:** REST, gRPC, GraphQL
- **Asynchronous:** Kafka, RabbitMQ, JMS, AWS SQS

#### **Q10. REST vs Messaging vs gRPC**

<b>Feature</b>	<b>REST</b>	<b>gRPC</b>	<b>Messaging</b>
Type	Request-Response	Request-Response	Async
Format	JSON	Protobuf	Custom
Best for APIs		High perf comm	Decoupled systems

#### **Q11. What is service discovery?**

- Automatically find service locations instead of hardcoding URLs.  
Example: **Eureka, Consul, Kubernetes service registry**

#### **Q12. How to achieve load balancing?**

- Client-side (Spring Cloud LoadBalancer)
- Server-side (NGINX, K8s Services)
- API Gateway level

### **Data Management & Consistency**

#### **Q13. How do microservices handle data?**

- Each service **owns its database** (no shared schema).
- Communication via APIs or events (event sourcing).

#### **Q14. What is the Saga Pattern?**

- A way to manage distributed transactions.
- Each service executes a local transaction and publishes an event to trigger the next step.
- If one fails → compensating transactions undo previous actions.

#### **Q15. Event-driven architecture advantages?**

- Loose coupling
- Scalability
- Better resilience
- Easier async workflows

#### **Q16. CAP theorem**

- Consistency, Availability, Partition Tolerance — you can guarantee only 2 at a time.
- 

### **5 Resilience & Fault Tolerance**

#### **Q17. How do you make microservices resilient?**

- Circuit Breakers (Resilience4j)
- Bulkhead pattern
- Retry + Timeout
- Fallbacks
- Rate Limiting
- Centralized logging + tracing

#### **Q18. Explain Circuit Breaker pattern**

- Detects failure and stops requests to a failing service for a cool-off period.

#### **Q19. Difference between Retry and Circuit Breaker?**

- Retry: attempt again after failure.
  - Circuit Breaker: stops trying after repeated failures to avoid overloading.
- 

### **6 Security in Microservices**

#### **Q20. How do you secure microservices?**

- OAuth2 + JWT (token-based auth)
- API Gateway for token validation
- HTTPS communication
- Role-based access control (RBAC)

#### **Q21. How does token validation work?**

- Client → Auth Server → receives JWT
- JWT sent with every request
- Gateway or service validates JWT before allowing access

#### **Q22. What is the difference between OAuth2 and OpenID Connect?**

- OAuth2 handles authorization.
  - OpenID Connect adds authentication (user identity).
- 

### **7 Deployment & DevOps**

#### **Q23. How are microservices deployed?**

- Using Docker containers, Kubernetes, or cloud-native services (AWS ECS, GKE, Azure AKS).

#### **Q24. What are advantages of containerization?**

- Portability, scalability, faster deployment, resource isolation.

#### **Q25. What are Kubernetes features used for microservices?**

- Pods, Deployments, Services, ConfigMaps, Secrets, Ingress, Horizontal Pod Autoscaling (HPA).
- 

### 8 Monitoring, Logging, and Tracing

#### **Q26. How do you monitor microservices?**

- Metrics: Prometheus, Grafana
- Tracing: Sleuth, Zipkin, OpenTelemetry
- Logging: ELK (Elasticsearch, Logstash, Kibana)

#### **Q27. What is distributed tracing?**

- Tracks requests across multiple services using trace IDs and span IDs.
- 

### API Gateway and Edge Services

#### **Q28. What is the role of an API Gateway?**

- Entry point for all client requests.
- Handles authentication, routing, rate limiting, and load balancing.

#### **Q29. Why use Spring Cloud Gateway instead of Zuul?**

- Non-blocking, reactive design.
  - Better performance and easier integration with modern Spring Boot versions.
- 

### 10 Testing Microservices

#### **Q30. How do you test microservices?**

- **Unit Testing:** JUnit, Mockito
  - **Integration Testing:** TestContainers, WireMock
  - **Contract Testing:** Spring Cloud Contract
  - **End-to-End Testing:** Postman, Karate, Cucumber
- 

### 1 Advanced Scenario / Real-World Questions

#### **Q31. How to handle versioning in microservices?**

- URI versioning: /api/v1/orders
- Header-based versioning
- Backward compatibility for old clients

#### **Q32. What happens if one microservice is down?**

- Circuit Breaker trips → fallback response
- Gateway returns cached data or graceful error

#### **Q33. How do you manage inter-service dependencies?**

- Through asynchronous messaging or domain events.
- Avoid tight coupling.

#### **Q34. How do you handle distributed logging?**

- Use correlation IDs or trace IDs added by Sleuth/OpenTelemetry.

#### **Q35. How to migrate from monolith to microservices?**

- Use **Strangler Fig pattern**: Extract parts of monolith gradually into independent services while keeping the rest functional.

## 1 2 Experience-Based Questions (Asked in Interviews)

**Q36. Tell us about a microservices project you built.**

- Explain services, tech stack, patterns (Eureka, Gateway, Kafka, Resilience4j, Config Server).
- Mention challenges like config management or tracing and how you solved them.

**Q37. What was the most difficult part of building microservices?**

- Data consistency / Distributed transactions / Debugging multi-service issues.

**Q38. How did you monitor or handle failures in production?**

- Centralized monitoring, alerting, trace dashboards.

## BONUS: 2025 Trends in Microservices

Old	New (2025)
Hystrix	Resilience4j
Ribbon	Spring Cloud LoadBalancer
Zuul	Spring Cloud Gateway
Sleuth	OpenTelemetry
Monolithic deploy	K8s / Docker-based
REST only	REST + gRPC + Kafka

## Why do we need all these (Prometheus, Grafana, Zipkin/Sleuth, ELK) in Microservices?

Because in **microservices**, you don't have *one app* — you have *many small apps*, each running in its own container, pod, or VM.

When something fails, you must **see what happened across multiple services** at once — and that's impossible without **monitoring, logging, and tracing tools**.

## The 3 Pillars of Microservices Observability

Pillar	Tool	Purpose	Example
Metrics	 Prometheus + Grafana	Measure numeric data about system health	CPU usage, request rate, latency, memory
Logs	 ELK Stack (Elasticsearch, Logstash, Kibana)	Record detailed event logs	"OrderService failed for user X"
Traces	 Sleuth + Zipkin / OpenTelemetry	Follow a single request across multiple services	TraceID: "12345" travels Order → Payment → Inventory

## 1 Metrics (Prometheus + Grafana)

**Why needed:**

- To monitor **health and performance** of all services.

- Helps detect issues **before** users notice them.

### Prometheus

- Collects metrics (numbers) like:
  - http\_server\_requests\_seconds\_count
  - jvm\_memory\_used\_bytes
  - Custom metrics via @Timed, @Counted
- Pulls data from /actuator/prometheus endpoint in each service.

### Grafana

- Visualizes metrics with dashboards.
- You can see CPU, memory, latency trends, alerts, etc.

#### Example:

Prometheus collects metrics → Grafana shows charts like:

“Average response time of Payment Service over last 1h = 800ms”

---

## 2 Logging (ELK Stack)

### Why needed:

- Each service produces its own logs (hundreds of files).
- You can't SSH into each one — so you need a **centralized log system**.

### ELK Components:

- **Elasticsearch** → stores logs (searchable)
- **Logstash** → collects and parses logs
- **Kibana** → web UI to view logs and dashboards

#### Example:

You can search:

"ERROR" AND "OrderService" AND "traceId=abc123"

→ and instantly see all related logs across 10 services.

Without ELK, you'd spend hours tailing individual log files — with ELK, 5 seconds.

---

## 3 Tracing (Sleuth + Zipkin / OpenTelemetry)

### Why needed:

- When one user request passes through 5 services, you need to **trace** the entire flow.
- If the response was slow, tracing shows **where exactly it slowed down**.

### Spring Cloud Sleuth

- Adds a traceId and spanId to every log automatically.
- You can track how a request moves through different services.

### Zipkin / OpenTelemetry

- Collects and visualizes those traces.
- Displays which service took how long in a timeline view.

#### Example:

TraceID = abc123:

- API Gateway → 10ms
- Order Service → 30ms

- Payment Service → 500ms ⏳  
→ Root cause: Payment Service delay.
- 

### 🧠 Putting It All Together (Full Observability Stack)

Type	Tools	What You See	Why It's Important
Metrics	Prometheus + Grafana	CPU, latency, error rates	Detect trends, set alerts
Logs	ELK Stack	Detailed service logs	Debug exact issues
Traces	Sleuth + Zipkin/OpenTelemetry	End-to-end request flow	Identify bottlenecks, dependencies

---

### ✳️ Architecture Diagram (Conceptually)

```
[Microservice 1] \
[Microservice 2] --> [Actuator Metrics] --> Prometheus --> Grafana
[Microservice 3] /
|
|--> [Logs] --> Logstash --> Elasticsearch --> Kibana
|
|--> [Trace Data] --> Zipkin / OpenTelemetry Collector
```

---

### 🗣 In Interview, You Can Say:

“In our microservices setup, we use the full observability stack. Prometheus and Grafana for metrics and alerts, ELK stack for centralized logs, and Sleuth + Zipkin for distributed tracing. This combination helps us identify performance issues, debug production errors, and analyze request flows across services efficiently.”

a