

Spring Boot Advanced Interview Questions & Answers

Question 1: How would you handle inter-service communication in a microservice architecture using Spring Boot?

Answer:

In a microservice architecture, services often need to communicate with each other. The communication can be **synchronous** or **asynchronous**, depending on the use case.

1. Synchronous Communication (Direct Communication):

- For simple, direct communication between services, I would use **RestTemplate**.
- RestTemplate allows services to **send HTTP requests** and **receive responses**, similar to a two-way conversation between microservices.
- Example:
- ```
RestTemplate restTemplate = new RestTemplate();
```
- ```
String response = restTemplate.getForObject("http://service-name/api/data", String.class);
```

2. Declarative HTTP Client (Feign Client):

- For more **complex or large-scale interactions**, I would use **OpenFeign**.
- Feign Client simplifies service-to-service communication by allowing us to **declare REST clients using interfaces**, making the code **cleaner and more maintainable**.
- Example:
- ```
@FeignClient(name = "order-service")
```
- ```
public interface OrderClient {
```
- ```
 @GetMapping("/orders/{id}")
```
- ```
    Order getOrderById(@PathVariable Long id);
```
- ```
}
```

#### **3. Asynchronous Communication (Event-Driven):**

- For asynchronous communication—where immediate responses are not required—I would use **message brokers** such as **RabbitMQ** or **Apache Kafka**.
- These act like **community boards**, where one service can post messages (events) that other services can consume and act upon later.
- This ensures a **robust, decoupled, and fault-tolerant** communication system.

#### **Summary:**

 Use **RestTemplate** or **Feign Client** for **synchronous calls**.

 Use **Kafka** or **RabbitMQ** for **asynchronous event-driven communication**.

---

**Question 2: Can you explain the caching mechanism available in Spring Cache (Spring Boot)?****Answer:**

Caching in Spring Boot is like having a **memory box** where we temporarily store frequently used data so that it can be retrieved quickly without repeating expensive operations.

- The **Spring Cache abstraction** provides a simple and consistent way to integrate various caching solutions (like Ehcache, Hazelcast, Redis, etc.) into a Spring application.
- It acts as a **smart memory layer** for our operations, designed to **save time and resources** by storing the results of expensive computations—such as database queries or API calls.
- When the same data is requested again, Spring retrieves it **directly from the cache** instead of performing the full operation again.

**Key Benefits of Caching:**

1. Reduces database load and improves performance.
2. Enhances response time for frequently accessed data.
3. Minimizes redundant computations or external API calls.

**Common Cache Providers:**

- ConcurrentHashMap (default, simple in-memory cache)
- Ehcache
- Hazelcast
- Redis (commonly used in distributed systems)

---

**Question 3: How would you implement caching in a Spring Boot application?****Answer:**

To implement caching in a Spring Boot application, follow these steps:

1. **Add Cache Dependency:**
  - Include the required cache dependency in your pom.xml or build.gradle.
  - Example for Maven:

```
<dependency>
 <groupId>org.springframework.boot</groupId>
 <artifactId>spring-boot-starter-cache</artifactId>
</dependency>
```
  - Add provider-specific dependencies (like Redis, Ehcache, or Hazelcast) if needed.
2. **Enable Caching:**
  - Add the @EnableCaching annotation to your main application class.
  - @SpringBootApplication

- @EnableCaching
- public class Application {
- public static void main(String[] args) {
- SpringApplication.run(Application.class, args);
- }
- }

### 3. Define Cacheable Operations:

- Use the @Cacheable annotation on methods whose results should be cached.
- @Cacheable("users")
- public User getUserById(Long id) {
- return userRepository.findById(id).orElse(null);
- }

### 4. Manage Cache Updates and Eviction:

- Use @CachePut to update the cache when data changes.
- Use @CacheEvict to remove entries when they are no longer valid.
- @CacheEvict(value = "users", allEntries = true)
- public void clearCache() {
- // logic to clear cache
- }

### 5. Choose and Configure a Cache Provider:

- You can use the **default ConcurrentMapCacheManager** (in-memory).
- Or choose a more robust cache provider like **Ehcache**, **Redis**, or **Hazelcast** for distributed caching.
- Example (application.properties for Redis):
- spring.cache.type=redis
- spring.redis.host=localhost
- spring.redis.port=6379

### Summary:

👉 Add dependency → Enable caching → Use @Cacheable → Manage with @CachePut and @CacheEvict → Configure cache provider.

---

#### Extra Points to Mention in an Interview

- You can combine caching with **@Scheduled tasks** to refresh caches periodically.
- Always consider **cache invalidation strategies** to avoid stale data issues.
- For distributed systems, **Redis** or **Hazelcast** is recommended for better scalability.
- You can monitor cache performance using **Spring Boot Actuator** metrics.

## 1 Spring Boot Performance Issues

### Question:

Your Spring Boot application is experiencing performance issues under high load. What are the steps you would take to identify and address the performance problems?

### Answer:

First, I would identify the specific performance issues using **monitoring tools** such as **Spring Boot Actuator**, **Micrometer**, or **SPLUNK / Prometheus / Grafana** to get insights on memory usage, CPU utilization, and response times.

Next, I would **analyze application logs and metrics** to spot any **patterns or errors**, especially under high load conditions. Logs may show thread contention, slow database queries, or frequent garbage collection pauses.

Then, I would **conduct a performance test** to replicate the issue using tools like **JMeter** or **Gatling**. This helps simulate high user load and identify bottlenecks.

I would also use a **profiler** (such as VisualVM, JProfiler, or YourKit) for **code-level analysis** to identify slow methods, inefficient loops, or memory leaks.

After collecting findings, I would proceed to **optimize the code** and configuration:

- Optimize database queries by using indexes and reducing N+1 query problems.
- Implement **caching** using tools like **Spring Cache**, **Redis**, or **Caffeine** to reduce repetitive calls.
- Enable **connection pooling** (e.g., HikariCP tuning).
- Check for unnecessary synchronization or blocking operations in code.
- Review **thread pool configurations** in web servers or async executors.
- Enable **lazy loading** where appropriate to reduce data load size.
- Use **asynchronous processing** for time-consuming tasks.
- Evaluate **GC (Garbage Collector) tuning** if memory pressure is high.
- Consider **scaling** options – horizontal (more instances) or vertical (larger servers).

Finally, it's crucial to **continuously monitor the application** after deployment to ensure no further degradation occurs and to set **alerts** for key metrics like response time, throughput, and error rate.

---

## 2 Best Practices for Versioning REST APIs

### Question:

What are the best practices for versioning REST APIs in a Spring Boot application?

### Answer:

For versioning REST APIs in Spring Boot, the best practices include the following approaches:

#### 1. URL Versioning (most common)

- Include the version number directly in the URL path.

Example:

- /api/v1/products
- /api/v2/products
- Pros: Clear, easy to use, browser-friendly.
- Cons: May lead to many duplicated endpoints over time.

## 2. Header Versioning

- Specify the version in a **custom header** (e.g., X-API-VERSION: 1).
- Example:
- GET /api/products
- Header: X-API-VERSION=2
- Pros: Cleaner URLs, no breaking changes in routes.
- Cons: Harder to test manually without special tools.

## 3. Media Type Versioning (a.k.a. Content Negotiation or Accept Header Versioning)

- Define version in the Accept header.
- Example:
- Accept: application/vnd.company.app-v1+json
- Pros: Adheres to RESTful principles, flexible.
- Cons: Complex to implement and maintain.

## 4. Parameter Versioning

- Specify version as a **query parameter**.
- Example:
- GET /api/products?version=1
- Pros: Easy to test and implement.
- Cons: Less RESTful and less discoverable.

### Additional Best Practices:

- Always **document versioning strategy** in your API documentation (e.g., Swagger or OpenAPI).
- **Deprecate old versions gracefully**, giving clients enough notice.
- Use **semantic versioning** where possible.
- Maintain **backward compatibility** as long as feasible.

## 3 How Does Spring Boot Simplify the Data Access Layer?

### Question:

How does Spring Boot simplify the data access layer implementation?

### Answer:

Spring Boot simplifies the implementation of the data access layer by offering several **streamlined features** and **auto-configurations**:

1. It **auto-configures essential settings** such as the **DataSource**, **JPA**, and **Hibernate** configurations based on the dependencies present in the classpath — reducing manual setup.
2. It provides **built-in repository support**, such as **JpaRepository**, **CrudRepository**, and **PagingAndSortingRepository**, enabling easy CRUD operations without writing boilerplate code.
3. Spring Boot can automatically **initialize database schemas and seed data** using SQL or schema scripts placed in resources (**schema.sql**, **data.sql**).
4. It integrates **smoothly with various databases and ORM technologies**, including JPA, JDBC, MongoDB, and R2DBC.
5. It automatically **translates SQL exceptions** into Spring's consistent **DataAccessException** hierarchy, providing uniform and simplified error handling across different database types.
6. It supports **transaction management** via simple annotations like **@Transactional**, reducing complexity in managing database transactions.
7. It allows **profiles and environment-based configuration**, so you can easily switch between dev, test, and prod databases.
8. It enables **connection pooling** through HikariCP by default, optimizing database performance.
9. It works well with **Spring Data JPA's derived queries** and **custom JPQL/native queries**, giving developers flexibility.

**In short:**

Spring Boot reduces boilerplate, provides smart defaults, integrates with ORM frameworks seamlessly, and ensures consistency — making the data access layer **efficient, clean, and developer-friendly**.

## 1 What are Conditional Annotations?

**Question:**

What are conditional annotations, and explain the purpose of conditional annotations in Spring Boot.

**Answer:**

Conditional annotations in Spring Boot help us **create beans or configurations only if certain conditions are met**.

They act like **rules or switches** — “*If this condition is true, then load this bean or configuration.*”

A common example is **@ConditionalOnClass**, which creates a bean **only if a specific class is present on the classpath**. For example, Spring Boot auto-configures **DataSource** only if the **HikariDataSource** class exists in the classpath.

This makes our application **flexible, modular, and adaptable** to different environments **without changing code**.

In essence, conditional annotations:

- Improve modularity and efficiency.
- Enable context-aware configuration.
- Allow multiple environments (like dev, test, prod) to share the same code but load different configurations.

### Common Conditional Annotations:

- `@ConditionalOnClass` – activates config if a specific class is present.
- `@ConditionalOnMissingClass` – activates config if a class is absent.
- `@ConditionalOnBean` – loads if a certain bean exists.
- `@ConditionalOnMissingBean` – loads only if no bean of that type exists.
- `@ConditionalOnProperty` – based on property value in `application.properties`.
- `@ConditionalOnWebApplication` – loads only in web environments.
- `@ConditionalOnExpression` – uses SpEL expressions to determine condition.

These annotations make Spring Boot's **auto-configuration intelligent** and **environment-specific**.

---

## 2 Role of `@EnableAutoConfiguration` in Spring Boot

### Question:

Explain the role of `@EnableAutoConfiguration` annotation in a Spring Boot application.

How does Spring Boot achieve auto-configuration internally?

### Answer:

The `@EnableAutoConfiguration` annotation tells Spring Boot to **automatically configure the application based on the dependencies available** in the classpath.

Internally, Spring Boot uses **conditional evaluation** — it examines:

- The **classpath** (to see which libraries are present),
- The **existing beans** in the context,
- And the **properties** defined in `application.properties` or `application.yml`.

Based on these, it decides what configurations to apply.

Auto-configuration uses **conditional annotations** under the hood (like `@ConditionalOnClass`, `@ConditionalOnMissingBean`, etc.) inside its auto-configuration classes.

This smart mechanism allows Spring Boot to:

- Automatically configure components like `DataSource`, `JPA`, `WebMvc`, `Security`, etc.
- Avoid manual XML or Java configuration.
- Simplify and speed up development.

For example:

If Spring Boot detects **spring-boot-starter-web** on the classpath, it automatically configures **Tomcat**, **DispatcherServlet**, and **WebMvc** beans.

Thus, **@EnableAutoConfiguration** combined with conditional annotations is what makes Spring Boot applications **run out-of-the-box** with minimal setup.

---

### **3 Spring Boot Actuator Endpoints**

**Question:**

What are Spring Boot Actuator endpoints?

**Answer:**

Spring Boot Actuator is like a **toolbox for monitoring and managing** your Spring Boot application.

It provides **endpoints** that expose runtime information such as:

- **Health** (/actuator/health)
- **Metrics** (/actuator/metrics)
- **Environment** (/actuator/env)
- **Beans** (/actuator/beans)
- **Mappings, loggers, info**, and many others.

These endpoints help us **monitor, diagnose, and manage** the app in production — like checking health, performance, configurations, and system metrics.

However, since these endpoints can expose sensitive internal data, it's crucial to **secure them** properly.

---

### **4 How to Secure Actuator Endpoints**

**Question:**

How can we secure the actuator endpoints?

**Answer:**

There are several strategies to secure actuator endpoints:

#### **1. Limit Exposure**

- By default, not all actuator endpoints are exposed.
- You can control exposure using  
`management.endpoints.web.exposure.include` in `application.properties`.
- Example:  
`management.endpoints.web.exposure.include=health,info,metrics`

#### **2. Use Spring Security**

- Integrate Spring Security to **require authentication** for accessing actuator endpoints.
- Example:  
`http`

- .authorizeHttpRequests()
- .requestMatchers("/actuator/\*\*").hasRole("ACTUATOR\_ADMIN")
- .and()
- .httpBasic();

### 3. Use HTTPS Instead of HTTP

- Always prefer **HTTPS** to encrypt communication and protect sensitive metrics or configuration data.

### 4. Define Actuator Roles

- Create a **specific role**, e.g., ACTUATOR\_ADMIN, and assign it only to trusted users.
- This ensures that only authorized users can access internal application details.

These steps help ensure **observability without compromising security**.

---

## 5 Strategies for Spring Boot Performance Optimization

### Question:

What strategies would you use to optimize the performance of a Spring Boot application?

### Answer:

If a Spring Boot application is taking too long to respond, I would use the following strategies:

- **Implement caching** for frequently accessed data using Spring Cache with Redis or Caffeine.
- **Optimize database queries** — add indexes, use pagination, reduce N+1 selects.
- **Use asynchronous methods** (@Async) for non-blocking operations like sending emails or notifications.
- **Use a load balancer** (e.g., NGINX or AWS ELB) when traffic is high.
- **Review time complexity** of algorithms and code sections that handle large data.
- **Use reactive programming (Spring WebFlux)** to handle many concurrent requests efficiently.
- **Enable compression and connection pooling** for HTTP clients.
- **Profile and monitor** the app regularly to find bottlenecks (using JProfiler, Actuator, or Prometheus).
- **Use lightweight DTOs** for data transfer and avoid sending unnecessary payloads.

By combining these techniques, we can ensure high responsiveness and scalability even under heavy load.

---

## 6 Handling Multiple Beans of the Same Type

### Question:

How can we handle multiple beans of the same type in Spring?

### **Answer:**

When there are multiple beans of the same type in Spring, ambiguity arises during dependency injection.

We can handle this in two main ways:

#### **1. Using @Qualifier**

- Specify exactly which bean to inject.
- Example:
- `@Bean("mysqlDataSource")`
- `public DataSource mysqlDataSource() {...}`
- 
- `@Bean("oracleDataSource")`
- `public DataSource oracleDataSource() {...}`
- 
- `@Autowired`
- `@Qualifier("mysqlDataSource")`
- `private DataSource dataSource;`
- Here, Spring knows which bean to use.

#### **2. Using @Primary**

- Mark one bean as the **default choice**.
- Example:
- `@Primary`
- `@Bean`
- `public DataSource defaultDataSource() {...}`

If a qualifier is not provided, Spring automatically injects the bean marked as @Primary. This approach is cleaner when one bean should generally be used unless explicitly overridden.

---

#### **In summary:**

Conditional annotations make configuration intelligent.

`@EnableAutoConfiguration` drives Spring Boot's automation.

Actuator provides visibility.

Security ensures protection.

Performance tuning keeps the system responsive.

Qualifier and Primary maintain clarity in bean injection.



## 7 Best Practices for Managing Transactions in Spring Boot

### Question:

What are some best practices for managing transactions in Spring Boot applications?

### Answer:

In Spring Boot, **transactions** ensure **data integrity and consistency** — especially when multiple database operations must either all succeed or all fail together.

#### Key Best Practices:

##### 1. Use @Transactional Annotation Properly

- Apply it at the **service layer**, not the controller or repository.
- Example:
- `@Service`
- `public class OrderService {`
- `@Transactional`
- `public void placeOrder(Order order) {`
- `orderRepository.save(order);`
- `paymentService.processPayment(order);`
- `}`
- `}`

Here, if payment fails, both actions roll back automatically.

##### 2. Transaction Propagation

- Defines how transactions interact when a method calls another transactional method.
- Common propagation levels:
  - REQUIRED (default): joins the existing transaction or creates a new one if none exists.
  - REQUIRES\_NEW: always starts a new transaction.
  - SUPPORTS: runs within a transaction if one exists; otherwise, runs non-transactionally.
  - NESTED: runs in a nested transaction (supported only with specific DBs).

##### 3. Set the Correct Isolation Level

- Prevents concurrency issues (dirty reads, phantom reads, etc.).
- Example:
- `@Transactional(isolation = Isolation.SERIALIZABLE)`

Common levels:

- `READ_UNCOMMITTED`, `READ_COMMITTED`, `REPEATABLE_READ`, `SERIALIZABLE`

##### 4. Use Rollback Rules

- By default, Spring rolls back only on **unchecked exceptions** (**RuntimeException**).
- You can customize rollback behavior:
- `@Transactional(rollbackFor = Exception.class)`

## 5. Avoid Long Transactions

- Keep transactional boundaries small — only around the logic that must be atomic.
- This reduces lock contention and improves performance.

## 6. Never Call **@Transactional** Methods from Within the Same Class

- Because internal method calls bypass Spring's proxy, the transaction won't start properly.

## 7. Prefer Declarative Over Programmatic Transactions

- Declarative transactions (**@Transactional**) are cleaner and easier to maintain.
- Programmatic (using **TransactionTemplate**) should be used only in complex or dynamic scenarios.

## 8. Enable Transaction Management

- Spring Boot automatically enables it with **@EnableTransactionManagement**, but it's good practice to include it explicitly when customizing configs.

### Bonus Tip:

Monitor transaction duration and rollback count using **Spring Actuator Metrics** or **database logs**.

---

## 8 Exception Handling in Spring Boot

### Question:

How do you handle exceptions globally in a Spring Boot application?

### Answer:

Spring Boot offers powerful mechanisms for **centralized and consistent exception handling**.

### Recommended Approaches:

#### 1. Global Exception Handling with **@ControllerAdvice**

- Create a global class to handle exceptions from all controllers.
- Example:
- `@ControllerAdvice`
- `public class GlobalExceptionHandler {`
- `@ExceptionHandler(ResourceNotFoundException.class)`
- `public ResponseEntity<String>`
- `handleNotFound(ResourceNotFoundException ex) {`

- return ResponseEntity.status(HttpStatus.NOT\_FOUND).body(ex.getMessage());
- }
- 
- @ExceptionHandler(Exception.class)
- public ResponseEntity<String> handleGeneric(Exception ex) {
 ○ return ResponseEntity.status(HttpStatus.INTERNAL\_SERVER\_ERROR).body("Unexpected error occurred");
 }
- }
- Keeps error messages consistent across your application.

## 2. Use Custom Exception Classes

- Define domain-specific exceptions like UserNotFoundException, InvalidRequestException for clarity.

## 3. Use @ResponseStatus Annotation

- To directly map exceptions to HTTP status codes.
- @ResponseStatus(HttpStatus.BAD\_REQUEST)
- public class InvalidInputException extends RuntimeException {
 ○ public InvalidInputException(String message) {
 ○ super(message);
 }
 }

## 4. Include Detailed Error Response Object

- Return structured JSON with fields like timestamp, message, error code, and path.

## 5. Log Exceptions Properly

- Use logging frameworks like **SLF4J** or **Logback** for error tracking and debugging.

## Spring Boot Application Layers and Their Responsibilities

### Question:

Can you explain the layered architecture of a Spring Boot application?

### Answer:

A well-designed Spring Boot application follows **layered architecture**, separating concerns clearly:

#### 1. Controller Layer (Web Layer)

- Handles HTTP requests and responses.
- Converts data to and from JSON using Jackson.

- Example: `@RestController`
- 2. **Service Layer (Business Logic)**
  - Contains the core business rules and logic.
  - Annotated with `@Service`.
  - This is where we usually apply transactions.
- 3. **Repository Layer (Data Access)**
  - Interacts with the database using Spring Data JPA, JDBC, etc.
  - Annotated with `@Repository`.
  - Example:
    - `public interface UserRepository extends JpaRepository<User, Long> {}`
- 4. **Model Layer (Entity/DTO Layer)**
  - Represents domain objects or data transfer objects.
  - Annotated with `@Entity`, `@Table`, etc.

This structure improves **maintainability, testing, and scalability**.

---

## 10 Spring Boot Logging Best Practices

### Question:

What are the best practices for logging in Spring Boot?

### Answer:

1. **Use SLF4J with Logback (default)**
  - Always log through SLF4J (`LoggerFactory.getLogger()`).
  - Example:
    - `private static final Logger logger = LoggerFactory.getLogger(MyService.class);`
2. **Log at Appropriate Levels:**
  - TRACE: fine-grained information for debugging.
  - DEBUG: debug details for developers.
  - INFO: general application progress.
  - WARN: potential issues.
  - ERROR: actual errors that need attention.
3. **Avoid Logging Sensitive Data** (like passwords, tokens, PII).
4. **Use Log Patterns and Rolling Policies**
  - Configure in `application.properties`:
  - `logging.file.name=app.log`
  - `logging.pattern.console=%d{yyyy-MM-dd HH:mm:ss} - %msg%n`
5. **Centralized Log Management**
  - Use tools like **ELK Stack**, **Grafana Loki**, or **Splunk** for large-scale monitoring.
6. **Structured Logging**
  - Include correlation IDs to trace logs across microservices.

## Summary

Topic	Key Takeaways
<b>Transactions</b>	Use @Transactional, manage propagation & rollback, keep transactions short
<b>Exception Handling</b>	Use @ControllerAdvice, custom exceptions, and clear response formats
<b>Application Layers</b>	Controller → Service → Repository → Entity
<b>Logging</b>	Use SLF4J, log at proper levels, avoid sensitive data



## Spring Boot Interview Notes — Testing, YAML, Profiles, and AOP

---

### 1 Testing in Spring Boot

#### Question:

How do you approach testing in a Spring Boot application?

#### Answer:

Testing in Spring Boot ensures our application behaves correctly before deploying — just like testing every part of a rocket before launch .

Spring Boot simplifies testing by providing built-in tools and annotations that make it easy to test both individual components and the entire application.

### ✓ Types of Tests in Spring Boot

#### 1. Unit Testing

- Focus: Testing small, isolated pieces of code (usually methods or classes).
- Goal: Verify that logic works correctly without involving Spring context or external systems.
- Tools:
  - **JUnit 5** (test framework)
  - **Mockito** (mocking dependencies)
- Example:
- `@ExtendWith(MockitoExtension.class)`
- `public class UserServiceTest {`
- 
- `@Mock`
- `private UserRepository userRepository;`
- 
- `@InjectMocks`
- `private UserService userService;`

```
•
• @Test
• void testFindUserById() {
• when(userRepository.findById(1L))
• .thenReturn(Optional.of(new User(1L, "Alice")));
•
•
• User user = userService.getUserById(1L);
•
•
• assertEquals("Alice", user.getName());
• }
• }
```

 **Tip:** No Spring context is loaded in unit tests — they're fast and lightweight.

---

## 2. Integration Testing

- Focus: Tests how multiple parts of the application work together (controller, service, repository, etc.).
- Uses Spring's **ApplicationContext** to load the environment.
- Annotation:
- `@SpringBootTest`
- Example:
- `@SpringBootTest`
- `@AutoConfigureMockMvc`
- `public class UserControllerTest {`
- 
- `@Autowired`
- `private MockMvc mockMvc;`
- 
- `@Test`
- `void test GetUser() throws Exception {`
- `mockMvc.perform(get("/users/1"))`
- `.andExpect(status().isOk());`
- `}`
- }

 This simulates a real environment and ensures all layers communicate properly.

---

### Key Annotations

#### **@SpringBootTest**

- Tells Spring Boot to load the full application context.

- Used for **integration testing** where multiple layers are tested together.
- Example:
- `@SpringBootTest`
- `public class ApplicationTests {`
- `@Test`
- `void contextLoads() {}`
- }

### **@MockBean**

- Creates a **mock version** of a Spring bean and injects it into the context.
- Used to isolate the unit under test from its dependencies.
- Example:
- `@SpringBootTest`
- `public class OrderServiceTest {`
- 
- `@MockBean`
- `private PaymentService paymentService;`
- 
- `@Autowired`
- `private OrderService orderService;`
- 
- `@Test`
- `void testPlaceOrder() {`
- `when(paymentService.processPayment(any())).thenReturn("SUCCESS");`
- `String result = orderService.placeOrder(new Order());`
- `assertEquals("SUCCESS", result);`
- `}`
- }

 This prevents hitting real services or databases during testing.

## **2 YAML vs Properties Files**

### **Question:**

What advantages does YAML offer over properties files in Spring Boot? Are there any limitations?

### **Answer:**

Spring Boot supports both **.properties** and **.yaml** (YAML) for configuration. YAML provides a cleaner and more readable format, especially for complex configurations.

### **Advantages of YAML**

## 1. Hierarchical Structure

Easier to define nested configurations:

2. server:
3. port: 8080
4. servlet:
5. context-path: /app

## 6. Readability

YAML's indentation makes configuration more organized and visually clear.

## 7. Supports Lists and Maps

8. mail:
9. recipients:
10. - admin@example.com
11. - user@example.com

## 12. Allows Comments

13. # This is the server configuration
14. server:
15. port: 8080

### ⚠ Limitations

- **Indentation-sensitive** → incorrect spacing can cause parsing errors.
- **Less familiar** for developers used to .properties format.
- **Difficult to manage** for very large or dynamic configurations.

### ✓ When to Use YAML:

- Ideal for **microservices** or projects with structured configuration.

---

## 3 Spring Boot Profiles

### Question:

Explain how Spring Boot profiles work and why they are useful.

### Answer:

**Spring Profiles** allow you to define separate configurations for different environments (like *dev*, *test*, and *prod*).

Think of profiles as **modes** — each mode activates different settings depending on where your application runs.

### ✓ How Profiles Work

- Define multiple configuration files:
  - `application.yml` → common configs
  - `application-dev.yml` → dev environment
  - `application-prod.yml` → production environment
- Example:

- # application-dev.yml
- server:
- port: 8081
- spring:
- datasource:
- url: jdbc:mysql://localhost:3306/dev\_db
- Activate a profile:
- spring.profiles.active=dev

or via command line:

java -jar app.jar --spring.profiles.active=prod

### Why Use Profiles

- Isolate environment-specific settings.
- Simplify deployment without code changes.
- Improve maintainability and flexibility.

 Example:

Different database URLs or logging levels for development vs. production.

---

## Aspect-Oriented Programming (AOP) in Spring

### Question:

What is Aspect-Oriented Programming (AOP) in Spring Framework?

### Answer:

Aspect-Oriented Programming (AOP) is a programming paradigm that helps **separate cross-cutting concerns** — like logging, transactions, and security — from business logic.

### Key Concepts

Term	Description
------	-------------

**Aspect** A module containing cross-cutting logic (e.g., logging, auditing).

**Advice** The action taken by an aspect (e.g., code that runs before/after a method).

**Join Point** A point in program execution (like a method call) where advice can be applied.

**Pointcut** Expression that selects specific join points.

**Weaving** Linking aspects with application code at runtime.

### Example

@Aspect

@Component

public class LoggingAspect {

```
 @Before("execution(* com.example.service.*(..))")
```

```
 public void logBefore(JoinPoint joinPoint) {
```

```
 System.out.println("Method called: " + joinPoint.getSignature().getName());
 }
}
```

### **Real-World Uses of AOP:**

- Logging method execution
- Measuring performance
- Implementing security checks
- Managing transactions

AOP helps keep code **modular**, **clean**, and **maintainable** by keeping cross-cutting logic separate.

---

### **Summary Table**

#### **Topic**    **Key Idea**

**Testing** Use @SpringBootTest for integration, @MockBean for isolation

**YAML** More readable, supports hierarchy, but indentation-sensitive

**Profiles** Enable environment-specific configurations easily

**AOP** Handles cross-cutting concerns like logging, security, and transactions

### **Spring Boot — Advanced Interview Notes (Set 4)**

**Topics:** Aspect-Oriented Programming (AOP), Spring Cloud, Embedded Server, Bean Listing, DI, and Performance Optimization

---

## **1 Aspect-Oriented Programming (AOP)**

### **Question:**

What is Aspect-Oriented Programming (AOP) and how does it help in Spring?

### **Answer:**

Aspect-Oriented Programming (AOP) is a programming approach that helps **separate cross-cutting concerns** — functionalities that are required across multiple parts of an application.

In a traditional program, you might add logging, security checks, or transaction management in multiple methods. AOP allows you to **define such common logic once** and apply it wherever needed, automatically.

This keeps your core business logic clean and focused.

### **Example**

Imagine an e-commerce app:

- Each service method should log execution time and verify user access.
- Instead of adding logging and security code in every method:
- You define a **LoggingAspect** and **SecurityAspect** once.
- Then you specify where they should run (e.g., before or after a method call).

This reduces duplication and improves maintainability.

### Key Concepts

#### Concept Description

**Aspect** A module that encapsulates cross-cutting logic (e.g., LoggingAspect).

**Advice** The action taken by an aspect (before, after, around method execution).

**Join Point** A point in the program (e.g., method execution) where advice can be applied.

**Pointcut** Expression that defines *where* advice should apply.

**Weaving** Linking aspects with application code at runtime.

### Example Code

@Aspect

@Component

```
public class LoggingAspect {
```

```
 @Before("execution(* com.example.service.*(..))")
 public void logBefore(JoinPoint joinPoint) {
 System.out.println("Method Called: " + joinPoint.getSignature().getName());
 }
}
```

### Benefits

- Cleaner and modular code
- Easier maintenance
- Centralized management of common logic like logging, transactions, and security

---

## 2 Spring Cloud

### Question:

What is Spring Cloud and how is it useful for building microservices?

### Answer:

**Spring Cloud** is a suite of tools built on top of the Spring Framework to simplify the development of **microservice-based architectures**.

Think of Spring Cloud as the “manager” that ensures all your microservices — authentication, payments, inventory, orders — **work together smoothly**.

### Why We Need It

In a large system like an online shopping platform:

- One microservice handles user login,
- Another handles the cart,
- Another processes payments, and
- Another lists products.

Coordinating and managing all these microservices (their discovery, communication, and security) is challenging — and that's where Spring Cloud helps.

### Key Features

Feature	Description
<b>Service Discovery (Eureka)</b>	Services automatically register and find each other.
<b>API Gateway (Spring Cloud Gateway)</b>	Manages and routes requests between clients and services.
<b>Config Server</b>	Centralized configuration management for all microservices.
<b>Circuit Breaker (Resilience4j)</b>	Handles failures gracefully to prevent system crashes.
<b>Load Balancing (Ribbon / Spring Cloud LoadBalancer)</b>	Distributes requests among multiple service instances.
<b>Distributed Tracing (Sleuth + Zipkin)</b>	Tracks requests across services for debugging and performance monitoring.

#### In short:

Spring Cloud helps connect, secure, and manage microservices — making distributed systems more stable, scalable, and maintainable.

---

## 3 Spring Boot Server Selection

### Question:

How does Spring Boot decide which server to use?

### Answer:

Spring Boot automatically decides which embedded server to use based on **classpath dependencies**.

### How It Works

- If **Tomcat** is on the classpath → it becomes the default server.
- If **Jetty** or **Undertow** are included instead → Spring Boot configures them automatically.
- If no server dependency is found → Spring Boot defaults to **Tomcat** (included in `spring-boot-starter-web`).

### Example

```
<dependency>
 <groupId>org.springframework.boot</groupId>
 <artifactId>spring-boot-starter-web</artifactId> <!-- includes Tomcat -->
</dependency>
```

To switch to Jetty:

```
<dependency>
 <groupId>org.springframework.boot</groupId>
 <artifactId>spring-boot-starter-jetty</artifactId>
</dependency>
```

 This **auto-detection** simplifies setup — no need to manually configure or install a web server.

---

## Listing All Beans in a Spring Boot Application

### Question:

How can you get the list of all beans in your Spring Boot application?

### Answer:

You can access all beans managed by Spring's **ApplicationContext**.

### Code Example

```
import org.springframework.beans.factory.annotation.Autowired;
import org.springframework.context.ApplicationContext;
import org.springframework.stereotype.Component;
```

```
@Component
public class BeanLister {
```

```
 @Autowired
 private ApplicationContext context;
```

```
 public void listAllBeans() {
 String[] beans = context.getBeanDefinitionNames();
 for (String bean : beans) {
 System.out.println(bean);
 }
 }
}
```

 This is useful for debugging and understanding what components Spring Boot has created and managed in the context.

---

## 5 Improving Spring Boot Performance

### Question:

Describe a Spring Boot project where you significantly improved performance. What techniques did you use?

### Answer:

In one of my Spring Boot projects, I improved performance using the following techniques:

#### Techniques Used

1. **Database Optimization**
  - Enabled **connection pooling** for efficient database access.
  - Added **query caching** using **EHCache** to avoid repeated DB calls.
2. **Asynchronous Processing**
  - Used @Async for non-critical background tasks (like sending emails).
3. **HTTP Response Compression**
  - Enabled gzip compression in application properties to reduce payload size.
4. server.compression.enabled=true
5. **Stateless Sessions**
  - Configured Spring Security to use stateless JWT-based sessions, reducing memory usage.
6. **Monitoring**
  - Integrated **Spring Boot Actuator** for real-time health and performance metrics.
7. **Load Balancing and Scaling**
  - Deployed the app behind a load balancer to handle high user traffic.

#### Result:

Response times improved significantly, the app handled more concurrent users, and resource utilization was optimized.

---

## 6 Embedded Servlet Containers

### Question:

Explain the concept of Spring Boot's embedded servlet containers.

### Answer:

Spring Boot includes an **embedded web server** (like Tomcat, Jetty, or Undertow) within your application — so you don't need to deploy it separately on an external server.

#### Benefits

- **Run Anywhere:** The app runs via a simple java -jar command.
- **Simplified Deployment:** All dependencies and server are bundled into one executable JAR/WAR.

- **Faster Development:** Ideal for local testing and CI/CD pipelines.

### Example

Run directly:

`mvn spring-boot:run`

or

`java -jar myapp.jar`

 This “**self-contained server**” model saves setup time and makes Spring Boot apps portable and cloud-ready.

---

## 7 Dependency Injection (DI) Simplification in Spring Boot

### Question:

How does Spring Boot make dependency injection easier compared to traditional Spring?

### Answer:

In traditional Spring, you had to manually configure beans using **XML** or **explicit annotations**, which was time-consuming and error-prone.

Spring Boot simplifies this with **auto-configuration** and **component scanning**:

- It automatically discovers beans using `@ComponentScan`.
- It auto-configures required dependencies based on what's available in the classpath.

### Example

If you add `spring-boot-starter-data-jpa`, Spring Boot:

- Detects JPA dependencies,
- Creates a `DataSource` and `EntityManagerFactory`,
- Configures Hibernate — automatically!

 This allows developers to **focus on business logic**, not repetitive configuration.

---

## Summary Table

Topic	Key Idea
AOP	Separates cross-cutting concerns like logging/security
Spring Cloud	Simplifies microservice communication and management
Server Selection	Auto-detects Tomcat, Jetty, or Undertow
List Beans	Use <code>ApplicationContext.getBeanDefinitionNames()</code>
Performance	Optimize DB, caching, async, compression, monitoring
Embedded Server	Self-contained web server — no external deployment
DI Simplification	Auto-configures and scans beans automatically

## **How does Spring Boot simplify testing and deployment compared to traditional Spring?**

Spring Boot simplifies testing and deployment through its **auto-configuration, embedded servers, and dependency management** features.

In traditional Spring, developers had to manually configure beans, XML files, and servers like Tomcat or Jetty. Spring Boot removes these steps by automatically configuring components based on the dependencies on the classpath and by including **embedded servers**.

This allows developers to run applications simply using mvn spring-boot:run or java -jar app.jar, without setting up an external server.

For testing, Spring Boot integrates with **JUnit and Mockito**, and provides specialized annotations like:

- `@SpringBootTest` – for loading the full Spring context.
  - `@DataJpaTest`, `@WebMvcTest`, `@MockBean` – for focused testing of data, controllers, or services.
- This simplifies setup, reduces boilerplate, and allows fast testing in isolated environments.

---

## **Explain Spring Boot Actuator and its importance**

Spring Boot **Actuator** is a built-in module that provides **production-ready features** to monitor and manage applications. It exposes useful endpoints that give insight into the running application.

Examples include:

- `/actuator/health` – shows whether the app is running correctly.
- `/actuator/metrics` – provides performance data like memory usage, HTTP requests, and more.
- `/actuator/env` – shows environment properties and configurations.

These endpoints help **developers and DevOps teams** identify performance bottlenecks, detect failures early, and track system health.

Actuator is essential for maintaining **observability** in microservice environments, especially when integrated with monitoring tools like **Prometheus** or **Grafana**.

---

## **How does Spring Boot handle exception handling globally?**

Spring Boot provides several ways to handle exceptions globally:

1. **@ControllerAdvice + @ExceptionHandler**

Used to define global exception handlers across multiple controllers.

2. **ResponseEntityExceptionHandler**

Can be extended to customize standard Spring MVC error responses.

### 3. Custom error pages or JSON responses

Using the ErrorController interface to return structured JSON or HTML error pages.

This global approach ensures consistent error handling, better debugging, and cleaner controller logic.

---

### Explain how configuration properties are managed in Spring Boot

Spring Boot supports flexible configuration management using:

- application.properties or application.yml
- Environment variables
- Command-line arguments
- External configuration servers (like Spring Cloud Config)

Using @ConfigurationProperties or @Value, properties can be injected directly into beans.

This approach supports **profile-based configurations** (like application-dev.yml or application-prod.yml) that help tailor behavior per environment.

### How would you manage externalized configuration and secure secrets in Spring Boot?

Spring Boot allows **externalized configuration**, meaning configuration values can be kept **outside the application code**, so they can easily change between environments (like dev, test, or prod) without rebuilding the app.

You can manage configurations in several ways:

#### 1. Property or YAML files:

Store settings in application.properties or application.yml.

Profiles like application-dev.yml or application-prod.yml help manage environment-specific values.

#### 2. Environment variables and command-line arguments:

These override file-based configs when the app runs, making them ideal for containerized or cloud deployments.

#### 3. Spring Cloud Config Server:

Allows you to store configurations centrally (for multiple services) and fetch them dynamically at runtime.

#### 4. Secret Management Integration:

For sensitive data such as passwords, tokens, and API keys, integrate with tools like **HashiCorp Vault**, **AWS Secrets Manager**, or **Azure Key Vault**.

Spring Boot provides built-in integration for these systems using the spring-cloud-starter-vault-config dependency.

This setup ensures secrets aren't hardcoded in files or exposed in version control, enhancing both **security** and **Maintainability**.

---

## What are some common performance optimization techniques in Spring Boot?

To improve performance in Spring Boot applications, you can apply the following strategies:

### 1. Enable caching:

Use Spring's @Cacheable annotation and a caching library like **Ehcache**, **Redis**, or **Caffeine** to reduce repetitive database calls.

### 2. Database optimization:

Use connection pooling (e.g., HikariCP), optimize queries, and prefer batch operations when possible.

### 3. Async processing:

Use @Async for non-critical background tasks to free up the main thread.

### 4. HTTP compression and GZIP:

Enable response compression to reduce bandwidth.

### 5. Profile tuning:

Use Spring Boot Actuator and Micrometer metrics to identify slow components and memory usage patterns.

### 6. Stateless design:

Make services stateless where possible to ease horizontal scaling and load balancing.

These techniques help ensure scalability, faster response times, and efficient resource usage.

---

## Explain how you would monitor and log a Spring Boot application in production

Monitoring and logging are crucial for maintaining healthy applications.

Here's how you can do it in Spring Boot:

### 1. Spring Boot Actuator:

Exposes operational endpoints such as /actuator/health, /actuator/metrics, /actuator/loggers for real-time monitoring.

### 2. Micrometer integration:

Works with Actuator to export metrics to tools like **Prometheus**, **Grafana**, or **New Relic**.

### 3. Centralized logging:

Use frameworks like **Logback** or **SLF4J** for structured logging. Logs can be sent to systems like **ELK Stack (Elasticsearch, Logstash, Kibana)** or **Graylog**.

### 4. Alerts and dashboards:

Integrate metrics with alert systems to detect anomalies early.

## 5. Distributed tracing:

Tools like **Zipkin** or **Jaeger** can trace requests across multiple microservices, helping detect latency and bottlenecks.

---

## How would you implement graceful shutdown in Spring Boot?

Graceful shutdown ensures that when an application stops, ongoing requests are completed properly instead of being cut off mid-way.

### 1. Enable graceful shutdown:

In application.yml:

2. server:

3. shutdown: graceful

### 4. Configure timeout:

You can set how long Spring waits for running tasks to finish using:

5. spring.lifecycle.timeout-per-shutdown-phase: 30s

### 6. Use pre-destroy hooks:

Annotate cleanup methods with @PreDestroy to release resources like DB connections or threads.

### 7. Use SmartLifecycle or custom listeners:

To perform additional shutdown logic or notify other services before termination.

This ensures stability and prevents data loss during restarts or deployments.

## How Can Spring Boot Be Integrated with Docker for Deployment?

Spring Boot integrates very well with Docker — allowing you to package your app into a lightweight, portable container.

Here's the step-by-step process:

### 1. Build your Spring Boot JAR

Run:

2. mvn clean package

This generates a target/myapp.jar file.

### 3. Create a Dockerfile

The Dockerfile defines how your image will be built:

4. FROM openjdk:17-jdk-slim

5. COPY target/myapp.jar app.jar

6. ENTRYPOINT ["java", "-jar", "/app.jar"]

- o FROM — selects the Java runtime.

- o COPY — adds your jar into the image.

- o ENTRYPOINT — tells Docker how to start the app.

### 7. Build the image:

8. docker build -t my-springboot-app .
9. **Run the container:**
10. docker run -p 8080:8080 my-springboot-app

This approach makes your Spring Boot app **portable**, consistent across environments, and ideal for **Kubernetes or cloud deployment**.

---

## How Can Spring Boot Be Integrated with Kubernetes?

Spring Boot apps can be deployed and scaled using **Kubernetes (K8s)**.

Steps:

1. Create a **Docker image** (as shown above).
2. Write a **Kubernetes Deployment YAML** to define pods and replicas.
3. apiVersion: apps/v1
4. kind: Deployment
5. metadata:
6. name: springboot-app
7. spec:
8. replicas: 3
9. selector:
10. matchLabels:
11. app: springboot-app
12. template:
13. metadata:
14. labels:
15. app: springboot-app
16. spec:
17. containers:
18. - name: springboot-app
19. image: my-springboot-app:latest
20. ports:
21. - containerPort: 8080
22. Add a **Service YAML** to expose it:
23. apiVersion: v1
24. kind: Service
25. metadata:
26. name: springboot-service
27. spec:
28. type: LoadBalancer
29. selector:
30. app: springboot-app

31. ports:
32. - port: 80
33. targetPort: 8080
34. Apply both configurations:
35. kubectl apply -f deployment.yaml
36. kubectl apply -f service.yaml

This setup ensures **high availability, auto-scaling, and centralized management.**

---

## How Can Spring Boot Handle Cross-Origin Resource Sharing (CORS)?

CORS (Cross-Origin Resource Sharing) controls how frontend and backend interact across different domains.

You can handle CORS in Spring Boot using:

1. **Global configuration:**
2. @Configuration
3. public class MyConfig {
4.     @Bean
5.     public WebMvcConfigurer corsConfigurer() {
6.         return new WebMvcConfigurer() {
7.             @Override
8.             public void addCorsMappings(CorsRegistry registry) {
9.                 registry.addMapping("/\*\*")
10.                 .allowedOrigins("http://localhost:3000")
11.                 .allowedMethods("GET", "POST", "PUT", "DELETE");
12.             }
13.         };
14.     }
15. }

### 16. At Controller level:

17. @RestController
18. @CrossOrigin(origins = "http://localhost:3000")
19. public class MyController { ... }

### 20. With SecurityConfig:

21. http.cors().and().csrf().disable();

This setup allows your React, Angular, or Vue frontend to safely call Spring Boot APIs.

---

## How Can Spring Boot Work with Kafka?

Spring Boot integrates with **Apache Kafka** for real-time messaging between microservices.

Steps:

1. Add dependency:

```

2. <dependency>
3. <groupId>org.springframework.kafka</groupId>
4. <artifactId>spring-kafka</artifactId>
5. </dependency>
6. Configure properties:
7. spring:
8. kafka:
9. bootstrap-servers: localhost:9092
10. consumer:
11. group-id: my-group
12. Create a producer:
13. @Service
14. public class ProducerService {
15. @Autowired
16. private KafkaTemplate<String, String> kafkaTemplate;
17.
18. public void sendMessage(String message) {
19. kafkaTemplate.send("my_topic", message);
20. }
21. }
22. Create a consumer:
23. @Service
24. public class ConsumerService {
25. @KafkaListener(topics = "my_topic", groupId = "my-group")
26. public void listen(String message) {
27. System.out.println("Received: " + message);
28. }
29. }

```

This allows microservices to **communicate asynchronously**, improving scalability and decoupling.

---

## How Can You Handle Errors Gracefully in Spring Boot?

Error handling is essential for user-friendly APIs.

1. **Use @ControllerAdvice and @ExceptionHandler**
2. @ControllerAdvice
3. public class GlobalExceptionHandler {
- 4.
5. @ExceptionHandler(Exception.class)
6. public ResponseEntity<String> handleAll(Exception ex) {

```
7. return new ResponseEntity<>("Error: " + ex.getMessage(),
 HttpStatus.BAD_REQUEST);
8. }
9. }
```

## 10. Custom Error Responses

Create a structured response class (timestamp, message, status).

## 11. ErrorController

Override Spring Boot's default /error endpoint if you need custom HTML or JSON error messages.

## 12. Validation errors

Use @Valid and @RestControllerAdvice to catch  
MethodArgumentNotValidException.

This ensures clean and consistent responses across your APIs.

---

Would you like me to continue next with:

- ◆ **Spring Boot with Database (JPA, Transactions, etc.)**
- ◆ **Microservices Communication (Feign, RestTemplate, etc.)**
- ◆ **Spring Boot Advanced Security (OAuth2, JWT Refresh Token Flow)**

## Resilience in Spring Boot Applications

To make a Spring Boot application more resilient, especially in **microservices architectures**:

- **Circuit Breakers:** Stop calls to failing services to prevent cascading failures. (e.g., **Resilience4j**)
  - **Retry Logic:** Automatically retry failed operations for transient errors.
  - **Timeouts:** Avoid blocking resources waiting for unresponsive services.
  - **Monitoring & Logging:** Track system health to detect and fix issues quickly.
- 

## Serverless Functions with Spring Cloud

- **Spring Cloud Function** allows writing business logic as simple Java functions.
  - Functions can run serverlessly on cloud platforms, automatically scaling to demand.
  - Focus is on code; Spring Cloud handles deployment, scaling, and execution.
- 

## Spring Cloud Gateway

- **Routing:** Configure routes via properties or Java config to direct requests.
- **Security:** Integrate with Spring Security for authentication/authorization.

- **Monitoring:** Use Spring Boot Actuator for health checks and metrics.
- 

## Synchronous & Asynchronous Task Management

- **Synchronous Tasks:** Integrate with **RabbitMQ or Kafka**; use messaging queues for processing and monitoring.
  - **Asynchronous Tasks:**
    - Use `@Async` to run tasks in the background.
    - Track progress with `CompletableFuture`.
    - Manage threads with **ThreadPoolTaskExecutor**.
    - Monitor tasks with **Spring Boot Actuator** for metrics.
- 

## Spring Security Configuration

- Use **Spring Security dependency**.
  - Extend `WebSecurityConfigurerAdapter` (or `SecurityFilterChain` for newer versions).
  - Protect endpoints with `.authorizeRequests()` and configure roles.
  - Enable **form login** or JWT-based authentication.
  - Use `@ConditionalOnMissingBean` to allow auto-configuration to back off when a custom bean exists.
- 

## Deployment

- **JAR:** Use embedded servers (Tomcat/Jetty).
  - `mvn package`
  - `java -jar target/app.jar`
  - **WAR:** Extend `SpringBootServletInitializer` for deployment to external servlet containers.
- 

## Configuration Features

- **Relaxed Binding:** Supports multiple property name formats (`server.port`, `server_port`, `server-port`).
  - **Externalized Configuration:** Manage sensitive info via Spring Cloud Config, environment variables, or encrypted secrets.
- 

## Spring Boot & CI/CD

- Integrate with **Jenkins**, **GitHub Actions**, or similar tools for automated build, test, and deployment pipelines.
  - Ensures faster feedback, fewer manual errors, and consistent deployments.
- 

## Server Flexibility

- Replace embedded Tomcat with Jetty or Undertow by excluding Tomcat and including the desired server in dependencies.
- 

## Error Handling

- **White-label error page:** Occurs when URLs are unmapped. Fix by mapping controllers correctly.
  - **Custom error handling:** Use `@ControllerAdvice` or implement `ErrorController`.
  - **404 Handling:** Map errors via `@RequestMapping("/error")` and customize response.
- 

## Pagination

- Use **Spring Data JPA** Pageable interface.
  - Example:
  - `PageRequest pageRequest = PageRequest.of(page, size);`
  - `Page<User> users = userRepository.findAll(pageRequest);`
- 

## Event-Driven Architecture

- Use **Spring events:**
    - Create custom events by extending `ApplicationEvent`.
    - Publish events using `ApplicationEventPublisher`.
    - Listen with `@EventListener`.
  - Helps decouple components for notifications, logging, etc.
- 

## Basic Spring Boot Annotations

- `@SpringBootApplication` → Combines `@Configuration`, `@EnableAutoConfiguration`, and `@ComponentScan`.
  - `@RestController + @RequestMapping` → Define REST endpoints.
  - `@Service, @Repository` → Mark service/data layers.
  - `@Autowired` → Dependency injection.
- 

## Distributed Tracing

- Tools: **Spring Cloud Sleuth, Zipkin**
  - Assigns unique IDs to requests across microservices.
  - Visualize request flow, identify delays, and troubleshoot issues.
- 

## Cloud Storage Integration

- Use **AWS S3 SDK or Google Cloud Storage**.
  - Configure credentials and properties.
  - Create a service class to manage file upload/download/delete operations.
-

## Rate Limiting

- Libraries: **Bucket4J** or **Spring Cloud Gateway**.
  - Limit requests per user/time frame to prevent abuse.
- 

## Soft Delete

- Add a deleted Boolean or deletedAt timestamp.
  - Update column instead of deleting records.
  - Filter out deleted records in repository queries.
- 

## Non-blocking Reactive REST APIs

- Use **Spring WebFlux** with Mono/Flux.
  - Reactive repositories like **ReactiveCrudRepository**.
  - Non-blocking I/O allows high concurrency and efficient resource usage.
- 

 This covers **advanced Spring Boot topics** including **resilience, microservices, serverless functions, reactive programming, security, deployment, CI/CD, and cloud integration** — all structured for interview readiness.